

APPENDIX

(A)

```
1 ---
2 title: "Distributed Data Analysis"
3 author: "2012754"
4 date: "21/03/21"
5 output: html_notebook
6 ---
7
8 This is an [R Markdown](http://rmarkdown.rstudio.com)
  Notebook. When you execute code within the notebook, the
  results appear beneath the code.
9
10 Try executing this chunk by clicking the *Run* button
  within the chunk or by placing your cursor inside it and
  pressing *Ctrl+Shift+Enter*.
11
12 #HOUSING DATA
13 ```{r}
14 library(dplyr)
15 #import dataset and inspect file
16 housing <- read.csv("C:/Users/tochi/Desktop/Housing_Data_
  GreaterLondon_2018-20.csv", header = TRUE)
17 str(housing)
18 summary(housing)
19
20 #drop columns
21 housing_clean <- select(housing, -'saon', -'paon',
  -'estate_type', -'street', -'linked_data_uri', -'county',
  -'transaction_category', -'street', -'locality')
22 view(housing_clean)
23 #convert into correct date format
24 housing_clean$deed_date <-
  as.Date(housing_clean$deed_date, format = "%Y-%m-%d")
25
26 #convert into categorical
27 housing_clean$property_type <-
  as.factor(housing_clean$property_type)
28 housing_clean$new_build <-
  as.factor(housing_clean$new_build)
29
30 #change deed date to date prior to our join
31 names(housing_clean)[3] <- 'date'
32
33 #format to just month and year
34 housing_clean$date <- format(housing_clean$date, "%m/%Y")
35
36 ```
```

(B)


```
36 ^  
37 ▾ ##DATA JOINING STREET CRIME WITH HOUSING  
38  
39 ▾ ```{r}  
40 library(zoo)  
41 library(anytime)  
42  
43  
44 #read in the street data and inspect it  
45 street_crime <-  
  read.csv('C:/Users/tochi/Desktop/metropolitan-street-final.csv')  
46 str(street_crime)  
47  
48 #convert into categorical  
49 street_crime$Crime.type <-  
  as.factor(street_crime$Crime.type)  
50 street_crime$postcode <- as.factor(street_crime$postcode)  
51  
52 street_crime$Month <- anytime(street_crime$Month)  
53 street_crime$Month <- as.Date(street_crime$Month, format  
  = "%Y-%m-%d")  
54 street_crime$Month <- format(street_crime$Month, '%m/%Y')  
55  
56 names(street_crime)[2] <- 'date'  
57  
58 #inner join on postcode  
59 Crime_Housing <- merge(housing_clean, street_crime, by =  
  c("postcode", "date"), all.x = TRUE, all.y = TRUE)  
60 #create a csv file with our df  
61 write.csv(Crime_Housing,  
  'C:/Users/tochi/Desktop/Crime_Housing.csv', row.names =  
  FALSE)  
62 str(Crime_Housing)  
63 plot(Crime_Housing$price_paid, ylab = "price_paid")  
64 plot(Crime_Housing$Crime.type, ylab = "Crimetype")  
65  
66 ^
```

(C)


```
68 ^
69 ▾ #Removing Outliers and Duplicate Rows
70 ▾ ```{r}
71
72 str(Crime_Housing)
73 summary(Crime_Housing)
74 boxplot(Crime_Housing$price_paid)
75 Crime_Housing_boxplot <-
  boxplot(Crime_Housing$price_paid)
76
77
78 Crime_Housing_boxplot$out
79
80 min(Crime_Housing_boxplot$out)
81
82 Crime_Housing[Crime_Housing$price_paid >=
  min(Crime_Housing_boxplot$out), ]
83
84 #to remove duplicate rows
85 Crime_Housing_clean <- unique(Crime_Housing)
86 str(Crime_Housing_clean)
87
88 ##to visualise graphically
89 Crime.type.freq <- xtabs(~ Crimetype, data =
  Crime_Housing)
90 Crime.type.prop <- prop.table(Crime.type.freq)
91 Crime.type.prop
92
93 Price_paid.hist <- ggplot(data = Crime_Housing, aes(x =
  price_paid)) + geom_histogram(binwidth =10, color =
  "black", fill = "lightgrey") + xlab("Price_Paid") +
  ylab("Frequency")
94
95
96 ^ ```
```

(D)

```
97 #EXPLORATORY DATA ANALYSIS
98 ```{r}
99 install.packages("ggally")
100 library(ggally)
101 #install.packages("mosaicData")
102
103 Price_crime_rel <- Crime_Housing_clean[, c("Crimetype",
104 "price_paid")]
105 ggpairs(Price_crime_rel)
106
107
108
109 ^```
```



i `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



(E)

```
In [1]: import pandas as pd
import datetime, timedelta

import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
```

```
In [2]: housing_Crime_Data = pd.read_csv(r"C:\Users\tochi\Desktop\Crime_Housing.csv")
housing_Crime_Data.describe()
```

```
Out[2]:
```

	price_paid
count	1.344350e+05
mean	9.119420e+05
std	5.360636e+06
min	1.000000e+02
25%	3.170000e+05
50%	4.350000e+05
75%	6.450000e+05
max	2.550000e+08

```
In [3]: housing_Crime_Data.dtypes
```

```
Out[3]: Month          object
Crime.type            object
postcode              object
price_paid            int64
property_type         object
new_build             object
town                  object
district              object
dtype: object
```

```
In [4]: housing_Crime_Data.columns
```

(F)

```
In [4]: housing_Crime_Data.columns
```

```
Out[4]: Index(['Month', 'Crime.type', 'postcode', 'price_paid', 'property_type',  
             'new_build', 'town', 'district'],  
            dtype='object')
```

```
In [5]: housing_Crime_Data['Crime.type'] = pd.Categorical(housing_Crime_Data['Crime.type'] )  
housing_Crime_Data['postcode'] = pd.Categorical(housing_Crime_Data['postcode'] )  
housing_Crime_Data['Month'] = pd.Categorical(housing_Crime_Data['Month'])  
housing_Crime_Data['property_type'] = pd.Categorical(housing_Crime_Data['property_type'] )  
housing_Crime_Data['new_build'] = pd.Categorical(housing_Crime_Data['new_build'] )  
housing_Crime_Data['district'] = pd.Categorical(housing_Crime_Data['district'] )  
housing_Crime_Data['town'] = pd.Categorical(housing_Crime_Data.loc[:, 'town'])
```

```
In [6]: housing_Crime_Data.dtypes
```

```
Out[6]: Month          category  
Crime.type          category  
postcode            category  
price_paid          int64  
property_type       category  
new_build           category  
town                category  
district            category  
dtype: object
```

```
In [8]: Housing_CrimeData_Clean_Cat = housing_Crime_Data.select_dtypes(include=['category'])  
Housing_CrimeData_Clean_num = housing_Crime_Data.select_dtypes(exclude=['category'])  
X_encoded = pd.get_dummies(Housing_CrimeData_Clean_Cat)
```

```
In [9]: frames = [X_encoded, Housing_CrimeData_Clean_num]  
combo_enc = pd.concat(frames, axis = 1)
```

```
In [10]: Housing_CrimeData_Clean_Cat_freq = Housing_CrimeData_Clean_Cat.copy()  
for c in Housing_CrimeData_Clean_Cat_freq.columns.tolist():  
    Housing_CrimeData_Clean_Cat_freq[c] = Housing_CrimeData_Clean_Cat_freq.groupby(c).transform('count')/len(Housing_CrimeData_Clean_Cat_freq)
```

```
In [11]: frames_freq = [Housing_CrimeData_Clean_Cat_freq, Housing_CrimeData_Clean_num]  
combo_enc_freq = pd.concat(frames_freq, axis = 1)
```

(G)

```
In [8]: Housing_CrimeData_Clean_Cat = housing_Crime_Data.select_dtypes(include=['category'])
Housing_CrimeData_Clean_num = housing_Crime_Data.select_dtypes(exclude=['category'])
X_encoded = pd.get_dummies(Housing_CrimeData_Clean_Cat)

In [9]: frames = [X_encoded, Housing_CrimeData_Clean_num]
combo_enc = pd.concat(frames, axis = 1)

In [10]: Housing_CrimeData_Clean_Cat_freq = Housing_CrimeData_Clean_Cat.copy()
for c in Housing_CrimeData_Clean_Cat_freq.columns.to_list():
    Housing_CrimeData_Clean_Cat_freq[c] = Housing_CrimeData_Clean_Cat_freq.groupby(c).transform('count')/len(Housing_CrimeData_Clean_Cat_freq)

In [11]: frames_freq = [Housing_CrimeData_Clean_Cat_freq, Housing_CrimeData_Clean_num]
combo_enc_freq = pd.concat(frames_freq, axis = 1)
combo_enc_freq.shape

Out[11]: (134435, 8)

In [12]: combo_enc_freq.tail()
combo_enc_freq.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 134435 entries, 0 to 134434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Month            134435 non-null float64
1   Crime.type       134435 non-null float64
2   postcode         134435 non-null float64
3   property_type    134435 non-null float64
4   new_build        134435 non-null float64
5   town             134435 non-null float64
6   district         134435 non-null float64
7   price_paid       134435 non-null int64  
dtypes: float64(7), int64(1)
memory usage: 8.2 MB

In [13]: from sklearn.model_selection import train_test_split
Xtrain, Xtest, ytrain, ytest = train_test_split(combo_enc_freq.filter(items = ['Month', 'Crime.type', 'postcode', 'property_type', 'new_build', 'town', 'district']), combo_enc_freq['price_paid'], random_state=0)
```

(H)

```
In [14]: from sklearn.neural_network import MLPRegressor
```

```
In [15]: from sklearn.datasets import make_regression
```

```
In [43]: regr = MLPRegressor(random_state=1, max_iter=15000).fit(Xtrain, ytrain)
```

```
C:\Users\tochi\anaconda3\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:582: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (15000) reached and the optimization hasn't converged yet.
warnings.warn(
```

```
In [44]: regr.predict(Xtest)
```

```
Out[44]: array([ 672271.06936196, 3195774.85376639, 489987.51534799, ...,
        3268704.75471761, 279349.94606583, 555175.15235468])
```

```
In [45]: regr.score(Xtest, ytest)
```

```
Out[45]: 0.027650184831768065
```

```
In [51]: #R_Square and Adjusted R Square
import statsmodels.api as sm
X_addC = sm.add_constant(Xtest)
result = sm.OLS(ytest, X_addC).fit()
print(result.rsquared, result.rsquared_adj)
```

```
0.019493542853790125 0.019289276754566154
```

```
In [69]: from sklearn.metrics import mean_squared_error
import math
pred_y = regr.predict(Xtest)
print(mean_squared_error(ytest, pred_y))
print(math.sqrt(mean_squared_error(ytest, pred_y)))
```

```
30351495581054.434
5509219.14440281
```

```
In [67]: from sklearn.metrics import mean_absolute_error
print(mean_absolute_error(ytest, pred_y))
```

```
807765.0093318874
```


(I-1)

```
In [1]: !pip install spark
```

```
Collecting spark
  Downloading spark-0.2.1.tar.gz (41 kB)
    |██████████████████████████████████████| 41 kB 204 kB/s
Building wheels for collected packages: spark
  Building wheel for spark (setup.py) ... done
  Created wheel for spark: filename=spark-0.2.1-py3-none-any.whl size=58738 sha256=d2ef5ba9f0c493d0136c93f2823ea36a0b39720c88587af5f94bfa2849ac1357
  Stored in directory: /Users/sne2909/Library/Caches/pip/wheels/c5/19/ff/9b16f354528bc9698ec3286be7947ebbf1f839132553961d4
Successfully built spark
Installing collected packages: spark
Successfully installed spark-0.2.1
```

```
In [2]: !pip install imblearn
```

```
Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Collecting imbalanced-learn
  Downloading imbalanced_learn-0.8.0-py3-none-any.whl (206 kB)
    |██████████████████████████████████████| 206 kB 4.8 MB/s
Requirement already satisfied: joblib>=0.11 in /Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages (from imbalanced_learn->imblearn) (0.15.1)
Requirement already satisfied: scipy>=0.19.1 in /Users/sne2909/Library/Python/3.8/lib/python/site-packages (from imbalanced_learn->imblearn) (1.4.1)
Requirement already satisfied: numpy>=1.13.3 in /Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages (from imbalanced_learn->imblearn) (1.19.5)
Collecting scikit-learn>=0.24
  Downloading scikit_learn-0.24.1-cp38-cp38-macosx_10_13_x86_64.whl (7.2 MB)
    |██████████████████████████████████████| 7.2 MB 15.4 MB/s
Requirement already satisfied: threadpoolctl>=2.0.0 in /Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages (from scikit_learn>=0.24->imbalanced_learn->imblearn) (2.1.0)
Installing collected packages: scikit-learn, imbalanced-learn, imblearn
  Attempting uninstall: scikit-learn
    Found existing installation: scikit-learn 0.23.1
    Uninstalling scikit-learn-0.23.1:
      Successfully uninstalled scikit-learn-0.23.1
Successfully installed imbalanced-learn-0.8.0 imblearn-0.0 scikit-learn-0.24.1
```

```
In [3]: !pip install pyspark
```

```
Collecting pyspark
```

(I-2)

```
Uninstalling scikit-learn-0.23.1:  
  Successfully uninstalled scikit-learn-0.23.1  
Successfully installed imbalanced-learn-0.8.0 imblearn-0.0 scikit-learn-0.24.1
```

```
In [3]: !pip install pyspark
```

```
Collecting pyspark  
  Downloading pyspark-3.1.1.tar.gz (212.3 MB)  
    [████████████████████████████████████████] 212.3 MB 57 kB/s  
Collecting py4j==0.10.9  
  Downloading py4j-0.10.9-py2.py3-none-any.whl (198 kB)  
    [████████████████████████████████████████] 198 kB 71.7 MB/s  
Building wheels for collected packages: pyspark  
  Building wheel for pyspark (setup.py) ... done  
  Created wheel for pyspark: filename=pyspark-3.1.1-py2.py3-none-any.whl size=212767604 sha256=3e650e1c9e5e6e78d7daaa54fe39db83f4d79199faeb7529e3ec45b21f5325fa  
    Stored in directory: /Users/sne2909/Library/Caches/pip/wheels/b3/0e/81/264aee961e43b9f6ba9ec81c8c540d2d7dccc52c6b51cbf22  
Successfully built pyspark  
Installing collected packages: py4j, pyspark  
Successfully installed py4j-0.10.9 pyspark-3.1.1  
Note: you may need to restart the kernel to use updated packages.
```

```
In [4]: from pyspark.sql import SQLContext  
from pyspark.sql import DataFrameNaFunctions  
from pyspark.ml import Pipeline  
from pyspark.ml.classification import DecisionTreeClassifier  
from pyspark.ml.classification import LogisticRegression  
from pyspark.ml.feature import Binarizer  
from pyspark.ml.feature import OneHotEncoder, VectorAssembler, StringIndexer, VectorIndexer  
from pyspark.ml.classification import RandomForestClassifier  
from pyspark.sql.functions import avg  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline  
from pyspark.ml.evaluation import MulticlassClassificationEvaluator  
from imblearn.over_sampling import SMOTE  
from imblearn.combine import SMOTEENN  
from pyspark.mllib.evaluation import MulticlassMetrics  
from pyspark.ml.evaluation import BinaryClassificationEvaluator  
from sklearn.model_selection import train_test_split  
from collections import Counter
```

```
In [5]: from pyspark.context import SparkContext
```

(I-3)

```
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9 pyspark-3.1.1
Note: you may need to restart the kernel to use updated packages.
```

```
In [4]: from pyspark.sql import SQLContext
        from pyspark.sql import DataFrameNaFunctions
        from pyspark.ml import Pipeline
        from pyspark.ml.classification import DecisionTreeClassifier
        from pyspark.ml.classification import LogisticRegression
        from pyspark.ml.feature import Binarizer
        from pyspark.ml.feature import OneHotEncoder, VectorAssembler, StringIndexer, VectorIndexer
        from pyspark.ml.classification import RandomForestClassifier
        from pyspark.sql.functions import avg
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        %matplotlib inline
        from pyspark.ml.evaluation import MulticlassClassificationEvaluator
        from imblearn.over_sampling import SMOTE
        from imblearn.combine import SMOTEENN
        from pyspark.mllib.evaluation import MulticlassMetrics
        from pyspark.ml.evaluation import BinaryClassificationEvaluator
        from sklearn.model_selection import train_test_split
        from collections import Counter
```

```
In [5]: from pyspark.context import SparkContext
        from pyspark.sql.session import SparkSession
        sc = SparkContext()
        spark = SparkSession(sc)
```

```
In [8]: crime_housing = sqlContext.read.csv("../Crime_Housing.csv", header=True, inferSchema=True)
        crime_housing.columns
```

```
Out[8]: ['Month',
         'Crimetype',
         'postcode',
         'price_paid',
         'property_type',
         'new_build',
         'town',
         'district']
```

```
In [10]: crime_housing.printSchema()
```

(J)


```
town',  
'district']
```

```
In [10]: crime_housing.printSchema()
```

```
root  
|-- Month: string (nullable = true)  
|-- Crimetype: string (nullable = true)  
|-- postcode: string (nullable = true)  
|-- price_paid: integer (nullable = true)  
|-- property_type: string (nullable = true)  
|-- new_build: string (nullable = true)  
|-- town: string (nullable = true)  
|-- district: string (nullable = true)
```

```
In [25]: import seaborn as sns  
#TO VISUALISE 1D view of my dataset  
from scipy.stats import norm  
sns.distplot(dataframe_2[['price_paid']], fit=norm, kde=False)
```

```
Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe2015eeeb0>
```



```
In [28]: dataframe_1.groupby('Crimetype').count().sort_values('Month', ascending=False)
```

```
Out[28]:
```

	Month	postcode	property_type	new_build	town	district
Crimetype						

(K)

```
In [28]: dataframe_1.groupby('Crimetype').count().sort_values('Month', ascending=False)
```

```
Out[28]:
```

	Month	postcode	property_type	new_build	town	district
Crimetype						
Anti-social behaviour	28256	28256	28256	28256	28256	28256
Violence and sexual offences	23130	23130	23130	23130	23130	23130
Vehicle crime	12724	12724	12724	12724	12724	12724
Burglary	8759	8759	8759	8759	8759	8759
Other theft	8561	8561	8561	8561	8561	8561
Criminal damage and arson	5754	5754	5754	5754	5754	5754
Public order	4719	4719	4719	4719	4719	4719
Drugs	3706	3706	3706	3706	3706	3706
Shoplifting	2980	2980	2980	2980	2980	2980
Robbery	2977	2977	2977	2977	2977	2977
Theft from the person	2735	2735	2735	2735	2735	2735
Bicycle theft	1826	1826	1826	1826	1826	1826
Other crime	860	860	860	860	860	860
Possession of weapons	561	561	561	561	561	561

```
In [43]: # replace predictor variable with categorical labels instead... using pandas quantile cut we can do that
crime_housing_clean = crime_housing.toPandas().copy()

pd.qcut(crime_housing_clean['price_paid'], q=[0, .1, 0.25, .5, .75, .9, 1]).unique()
```

```
Out[43]: [(220000.0, 317000.0], (99.999, 220000.0], (317000.0, 435000.0], (435000.0, 645000.0], (645000.0, 1100000.0], (1100000.0, 2550000.0]]
Categories (6, interval[float64]): [(99.999, 220000.0] < (220000.0, 317000.0] < (317000.0, 435000.0] < (435000.0, 645000.0] < (645000.0, 1100000.0] < (1100000.0, 255000000.0]]
```

```
In [52]: crime_housing_clean['price_paid_categories'] = pd.qcut(crime_housing_clean['price_paid'], q=[0, .1, 0.25, .5, .75, .9, 1], labels=None)
# extract month from Month (date) column
crime_housing_clean['month_int'] = crime_housing_clean['Month'].apply(lambda x: pd.to_datetime(x).month)
```

(L-1)

```
In [53]: crime_housing_clean.head()
```

```
Out[53]:
```

	Month	Crimetype	postcode	price_paid	property_type	new_build	town	district	price_paid_categories	month_int
0	01/11/2019	Anti-social behaviour	RM6 5PJ	250000	T	N	ROMFORD	BARKING AND DAGENHAM	220k-317k	1
1	01/11/2019	Anti-social behaviour	RM6 5JJ	215000	F	N	ROMFORD	BARKING AND DAGENHAM	< \$220k	1
2	01/11/2019	Burglary	RM6 5JP	290000	T	N	ROMFORD	BARKING AND DAGENHAM	220k-317k	1
3	01/11/2019	Criminal damage and arson	RM6 5PJ	250000	T	N	ROMFORD	BARKING AND DAGENHAM	220k-317k	1
4	01/11/2019	Criminal damage and arson	RM6 5PJ	250000	T	N	ROMFORD	BARKING AND DAGENHAM	220k-317k	1

```
In [107]: crime_housing_clean.town.nunique()
```

```
Out[107]: 72
```

```
In [77]: crime_housing_clean.columns = crime_housing_clean.columns.str.lower()
cols = crime_housing_clean.columns
```

```
In [108]: sdf_crime_housing = spark.createDataFrame(crime_housing_clean)
# Encode categorical variables before test/train split
categoricalCols = ['crimetype', 'property_type', 'town', 'new_build', 'district']
stages = []

featureIndexers = [StringIndexer(inputCol=catCol, outputCol=catCol+'Index') for catCol in categoricalCols]
stages += featureIndexers

labelIndexer = [StringIndexer(inputCol=labelCol, outputCol='label') for labelCol in ['price_paid_categories']]
stages += labelIndexer

numericCols = ['month_int']

assemblerInputs = [col + 'Index' for col in categoricalCols] + numericCols
assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [assembler]

pipeline = Pipeline(stages=stages)
sdf = pipeline.fit(sdf_crime_housing).transform(sdf_crime_housing)

selectedCols = ['label', 'features'] + cols.tolist()
```

(L-2)

```
pipeline = Pipeline(stages=stages)
sdf = pipeline.fit(sdf_crime_housing).transform(sdf_crime_housing)

selectedCols = ['label', 'features'] + cols.tolist()
sdf = sdf.select(selectedCols)
sdf.printSchema()
```

```
root
|-- label: double (nullable = false)
|-- features: vector (nullable = true)
|-- month: string (nullable = true)
|-- crimetype: string (nullable = true)
|-- postcode: string (nullable = true)
|-- price_paid: long (nullable = true)
|-- property_type: string (nullable = true)
|-- new_build: string (nullable = true)
|-- town: string (nullable = true)
|-- district: string (nullable = true)
|-- price_paid_categories: string (nullable = true)
|-- month_int: long (nullable = true)
```

In [109]: sdf.show(5)

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|label|      features|    month|      crimetype|postcode|price_paid|property_type|new_build|  town|      dis
trict|price_paid_categories|month_int|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|  3.0|[0.0,1.0,1.0,0.0,...|01/11/2019|Anti-social behav...| RM6 5PJ|  250000|          T|    N|ROMFORD|BARKING AND DAG
ENHAM|      $220k - $317k|      1|
|  4.0|[0.0,0.0,1.0,0.0,...|01/11/2019|Anti-social behav...| RM6 5JJ|  215000|          F|    N|ROMFORD|BARKING AND DAG
ENHAM|      < $220k|      1|
|  3.0|[3.0,1.0,1.0,0.0,...|01/11/2019|      Burglary| RM6 5JP|  290000|          T|    N|ROMFORD|BARKING AND DAG
ENHAM|      $220k - $317k|      1|
|  3.0|[5.0,1.0,1.0,0.0,...|01/11/2019|Criminal damage a...| RM6 5PJ|  250000|          T|    N|ROMFORD|BARKING AND DAG
ENHAM|      $220k - $317k|      1|
|  3.0|[5.0,1.0,1.0,0.0,...|01/11/2019|Criminal damage a...| RM6 5PJ|  250000|          T|    N|ROMFORD|BARKING AND DAG
ENHAM|      $220k - $317k|      1|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

(M)

```
-- new_build: string (nullable = true)
-- town: string (nullable = true)
-- district: string (nullable = true)
-- price_paid_categories: string (nullable = true)
-- month_int: long (nullable = true)
```

In [109]: sdf.show(5)

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|label|   features|   month|   crimetype|postcode|price_paid|property_type|new_build|   town|   dis
+-----+-----+-----+-----+-----+-----+-----+-----+
| 3.0|[0.0,1.0,1.0,0.0,...|01/11/2019|Anti-social behav...| RM6 5PJ|   250000|           T|   N|ROMFORD|BARKING AND DAG
ENHAM|   $220k - $317k|           1|
| 4.0|[0.0,0.0,1.0,0.0,...|01/11/2019|Anti-social behav...| RM6 5JJ|   215000|           F|   N|ROMFORD|BARKING AND DAG
ENHAM|   < $220k|           1|
| 3.0|[3.0,1.0,1.0,0.0,...|01/11/2019|           Burglary| RM6 5JP|   290000|           T|   N|ROMFORD|BARKING AND DAG
ENHAM|   $220k - $317k|           1|
| 3.0|[5.0,1.0,1.0,0.0,...|01/11/2019|Criminal damage a...| RM6 5PJ|   250000|           T|   N|ROMFORD|BARKING AND DAG
ENHAM|   $220k - $317k|           1|
| 3.0|[5.0,1.0,1.0,0.0,...|01/11/2019|Criminal damage a...| RM6 5PJ|   250000|           T|   N|ROMFORD|BARKING AND DAG
ENHAM|   $220k - $317k|           1|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

In [164]: # split data in train & test

```
train, test = sdf.randomSplit([0.8, 0.2], seed=25)
print("Training Dataset Count: " + str(train.count()))
print("Test Dataset Count: " + str(test.count()))

Training Dataset Count: 107432
Test Dataset Count: 27003
```


(N)

```
In [165]: dt = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'label', maxDepth = 3, maxBins=100)
dtModel = dt.fit(train)
predictions = dtModel.transform(test)
predictions.select('crimetype', 'postcode', 'property_type', 'new_build', 'town', 'district', 'month_int', 'label', 'prediction').show(10)
```

crimetype	postcode	property_type	new_build	town	district	month_int	label	prediction
Violence and sexu...	SW16 2SN	F	N	LONDON	LAMBETH	1	0.0	0.0
Violence and sexu...	SW9 9UQ	F	N	LONDON	LAMBETH	1	0.0	0.0
Vehicle crime	SW8 2EU	F	N	LONDON	LAMBETH	1	0.0	0.0
Criminal damage a...	SW16 6JD	F	N	LONDON	LAMBETH	1	0.0	0.0
Robbery	SW16 5LJ	F	N	LONDON	LAMBETH	1	0.0	0.0
Bicycle theft	SE24 0NS	F	N	LONDON	LAMBETH	1	0.0	0.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW16 3LJ	F	Y	LONDON	LAMBETH	1	0.0	0.0

only showing top 10 rows

```
In [166]: # evaluate decision tree
evaluator = MulticlassClassificationEvaluator(predictionCol='prediction', labelCol='label', metricName='accuracy')
accuracy = evaluator.evaluate(predictions)

print("Test Error = %g " % (1.0 - accuracy))
print("Accuracy = %g " % accuracy)

Test Error = 0.65415
Accuracy = 0.34585
```

(O)

```
In [167]: from pyspark.ml.classification import RandomForestClassifier
rf = RandomForestClassifier(featuresCol = 'features', labelCol = 'label', maxBins=100)
rfModel = rf.fit(train)
predictions = rfModel.transform(test)
predictions.select('crimetype', 'postcode', 'property_type', 'new_build', 'town', 'district', 'month_int', 'label', 'prediction').show(10)
```

crimetype	postcode	property_type	new_build	town	district	month_int	label	prediction
Violence and sexu...	SW16 2SN	F	N	LONDON	LAMBETH	1	0.0	1.0
Violence and sexu...	SW9 9UQ	F	N	LONDON	LAMBETH	1	0.0	1.0
Vehicle crime	SW8 2EU	F	N	LONDON	LAMBETH	1	0.0	1.0
Criminal damage a...	SW16 6JD	F	N	LONDON	LAMBETH	1	0.0	1.0
Robbery	SW16 5LJ	F	N	LONDON	LAMBETH	1	0.0	1.0
Bicycle theft	SE24 0NS	F	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW16 3LJ	F	Y	LONDON	LAMBETH	1	0.0	1.0

only showing top 10 rows

```
In [168]: evaluator = MulticlassClassificationEvaluator(predictionCol='prediction', labelCol='label', metricName='accuracy')
accuracy = evaluator.evaluate(predictions)

print("Test Error = %g " % (1.0 - accuracy))
print("Accuracy = %g " % accuracy)

Test Error = 0.598193
Accuracy = 0.401807
```

(P)

```
In [169]: lr = LogisticRegression(featuresCol = 'features', labelCol = 'label', maxIter=10)
lrModel = lr.fit(train)
```

```
In [170]: predictions = lrModel.transform(test)
predictions.select('crimetype','postcode','property_type','new_build','town','district','month_int','label','prediction').show(10)
```

crimetype	postcode	property_type	new_build	town	district	month_int	label	prediction
Violence and sexu...	SW16 2SN	F	N	LONDON	LAMBETH	1	0.0	0.0
Violence and sexu...	SW9 9UQ	F	N	LONDON	LAMBETH	1	0.0	0.0
Vehicle crime	SW8 2EU	F	N	LONDON	LAMBETH	1	0.0	0.0
Criminal damage a...	SW16 6JD	F	N	LONDON	LAMBETH	1	0.0	0.0
Robbery	SW16 5LJ	F	N	LONDON	LAMBETH	1	0.0	0.0
Bicycle theft	SE24 0NS	F	N	LONDON	LAMBETH	1	0.0	0.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW9 8QH	T	N	LONDON	LAMBETH	1	0.0	1.0
Anti-social behav...	SW16 3LJ	F	Y	LONDON	LAMBETH	1	0.0	1.0

only showing top 10 rows

```
In [171]: evaluator = MulticlassClassificationEvaluator(predictionCol='prediction', labelCol='label', metricName='accuracy')
accuracy = evaluator.evaluate(predictions)

print("Test Error = %g " % (1.0 - accuracy))
print("Accuracy = %g " % accuracy)
```

Test Error = 0.714254
Accuracy = 0.285746

In []: